

# Learning in Neural Networks: From Intuition to Backpropagation

Sayan Chaki]

March 20, 2025

# Outline

- 1 Motivation: How Learning Happens
- 2 Intuitive Understanding of Learning
- 3 Gradient Descent

# The Mystery of Learning

- How do we humans learn?
  - Our brains contain approximately 86 billion neurons
  - Learning involves forming and strengthening connections
  - No explicit programming required
- We observe, make mistakes, and adapt
  - Example: Learning to ride a bicycle
  - Fall → Adjust → Try again → Improve
- Learning is an iterative process of improvement
- Can we enable machines to learn in similar ways?

# Learning in Biological vs. Artificial Systems

## Biological Learning

- Synaptic connections strengthen or weaken
- Neurons that fire together, wire together (Hebb's Law)
- Experience-dependent plasticity
- Reinforcement through neurotransmitters

## Artificial Neural Networks

- Weight parameters increase or decrease
- Connections adjusted based on error
- Data-dependent adaptation
- Reinforcement through error signals


### Key Insight

Both systems learn through adjustment of connections based on feedback.

# Learning Metaphors

## The Child Learning to Walk

- Starts with no knowledge of balance or coordination
- Tries to stand, falls down (error signal)
- Brain processes what went wrong
- Adjusts muscle control (parameter updates)
- With each attempt, gradually improves
- Eventually masters walking without conscious thought



## Neural Network Parallel

Initial random weights → Forward pass → Error calculation → Weight adjustment → Improved performance

# Learning Metaphors

## Finding the Valley in Fog

- You're on a mountain in thick fog
- Goal: Reach the lowest point (valley)
- Can only feel the slope under your feet
- You move in the direction of steepest descent
- Take small steps to avoid overshooting
- Each step brings you closer to the valley





Figure: Descent

## Network Parallel

The error landscape has many dimensions (one for each weight), and gradient descent helps navigate this landscape to find the minimum error.

## The Orchestra Conductor

- Orchestra plays a piece (forward pass)
- Conductor hears the overall sound (output)
- Identifies which instruments are out of tune (error attribution)
- Provides specific feedback to each musician (backpropagation)
- Musicians adjust their playing (weight updates)
- Overall performance improves



## Neural Network Parallel

Backpropagation distributes error corrections throughout the network, telling each weight how much it contributed to the error.

# Real-World Examples of Learning Systems

- **Computer Vision**

- Image classification (Is this a cat or a dog?)
- Object detection (Where is the pedestrian in this image?)
- Facial recognition (Who is this person?)

- **Natural Language Processing**

- Sentiment analysis (Is this review positive or negative?)
- Machine translation (English to French)
- Text generation (Complete this paragraph)


- **Reinforcement Learning**

- Game playing (Chess, Go, video games)
- Robotics (Learning to walk, grasp objects)
- Autonomous driving




# Neural Networks: A Brief Overview

- Inspired by biological neural networks (our brains)
- Composed of interconnected nodes (neurons)
- Organized in layers:
  - Input layer: Receives data
  - Hidden layer(s): Processes information
  - Output layer: Provides predictions
- Each connection has a weight parameter
- Neurons apply an activation function to their input



# What Does a Neural Network Need to Learn?


- A neural network is a collection of connections
- Each connection has a **weight** (a number)
- Learning means finding the right weights
- Goal: Find weights that make predictions match reality



**Example:** If  $y = w_1x_1 + w_2x_2$  and we want  $y$  to be 1 when  $x_1 = 2$  and  $x_2 = 1$ , we need:

$$w_1 \times 2 + w_2 \times 1 = 1 \quad (1)$$

# Learning as Weight Adjustment



**Example:** We want  $y$  to be 1 when  $x_1 = 1, x_2 = 0$


- Current weights:  $w_1 = 0.3, w_2 = 0.4, w_3 = 0.5$
- Forward pass:  
$$h = \sigma(w_1x_1 + w_2x_2) = \sigma(0.3 \times 1 + 0.4 \times 0) = \sigma(0.3) = 0.574$$
- Output:  $y = \sigma(w_3h) = \sigma(0.5 \times 0.574) = \sigma(0.287) = 0.571$
- Error:  $y$  should be 1, but it's 0.571
- Question: Which weights should we adjust, and by how much?

# A Concrete Example: XOR Problem

The XOR function:

$x_1$	$x_2$	$\text{XOR}(x_1, x_2)$
0	0	0
0	1	1
1	0	1
1	1	0

- XOR is not linearly separable
- Requires at least one hidden layer
- Network must learn appropriate weights




# How Learning Works - Intuitively

- 1 Present an input to the network (e.g.,  $x_1 = 1, x_2 = 0$ )
- 2 Calculate the output using current weights
- 3 Compare output to desired target (e.g.,  $y_{target} = 1$ )
- 4 Calculate the error (difference):  $E = \frac{1}{2}(y - y_{target})^2$
- 5 Adjust weights to reduce error
- 6 Repeat with many examples

## Key Questions:

- How much should we adjust each weight?
- In which direction should we adjust?
- How do we distribute the error across all weights?


# The Learning Loop



- This loop constitutes one training iteration
- Training typically requires thousands to millions of iterations
- Each iteration improves the weights slightly
- Eventually, the weights converge to good values


# The Learning Process: Visualization

## Linear Regression Example:



- Blue line: True relationship ( $y = 0.5x + 1$ )
- Black dots: Training data
- Red lines: Model predictions during training
  - Solid: Initial prediction
  - Dashed: After some training
  - Dotted: Final prediction

# Learning as Error Minimization



- Errors (dashed lines) are the differences between predictions and targets
- We want to minimize these errors
- Total error:  $E = \sum_i \frac{1}{2}(y_i - \hat{y}_i)^2$
- Learning means finding weights that minimize  $E$


## Key Insight

Learning is an optimization problem.

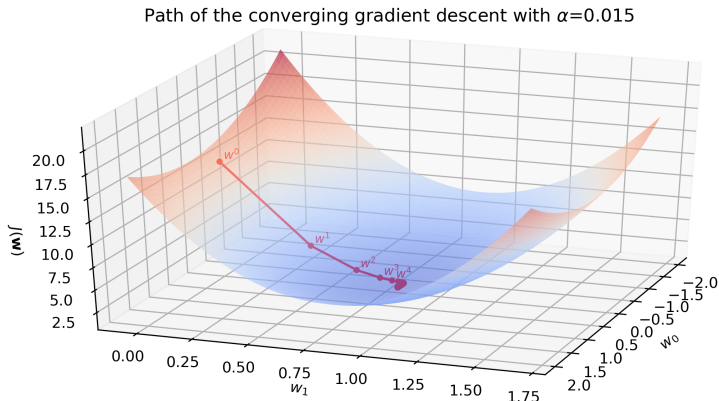


# The Error Landscape

- For each set of weights, there's an associated error
- This creates an "error landscape" over weight space
- Valleys represent low error (good predictions)
- Hills represent high error (poor predictions)
- Goal: Find the lowest point in this landscape




# Error Landscape: 2D Visualization



- x-axis: Weight 1
- y-axis: Weight 2
- z-axis (color): Error value
- Blue regions: Low error
- Red regions: High error

# Gradient Descent Intuition



- If we're on a slope, which way should we move?
- Answer: In the direction of steepest descent
- This direction is given by the negative gradient
- The gradient points "uphill"
- The negative gradient points "downhill"

## Key Insight

Follow the negative gradient to reduce error.

# What is a Gradient?

- The gradient is a vector of partial derivatives
- For a function  $f(x, y)$ , the gradient is:

$$\nabla f = \left( \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right) \quad (2)$$

- For error function  $E(w_1, w_2, \dots, w_n)$ :

$$\nabla E = \left( \frac{\partial E}{\partial w_1}, \frac{\partial E}{\partial w_2}, \dots, \frac{\partial E}{\partial w_n} \right) \quad (3)$$

- The gradient points in the direction of steepest increase
- The negative gradient points in the direction of steepest decrease

# The Gradient

The gradient of the error  $E$  with respect to a weight  $w$  tells us:


- Direction of steepest increase in error
- How sensitive the error is to changes in  $w$

Mathematically:  $\nabla E = \frac{\partial E}{\partial w}$

To minimize error, move in opposite direction:

$$w_{new} = w_{old} - \eta \frac{\partial E}{\partial w} \quad (4)$$

where  $\eta$  is the learning rate



# Simple Gradient Descent Example

**Linear regression with one weight:**

$$\hat{y} = wx \quad (5)$$

**Error function:**

$$E = \frac{1}{2}(y - \hat{y})^2 = \frac{1}{2}(y - wx)^2 \quad (6)$$

**Gradient:**

$$\frac{\partial E}{\partial w} = \frac{\partial}{\partial w} \frac{1}{2}(y - wx)^2 = -(y - wx)x = -xy + wx^2 \quad (7)$$


**Update rule:**

$$w_{new} = w_{old} - \eta(-xy + wx^2) = w_{old} + \eta(xy - wx^2) \quad (8)$$

**Interpretation**

If prediction is too low, increase  $w$ ; if too high, decrease  $w$ .

# Learning Rate



## Learning rate too small:

- Very slow convergence
- Many iterations needed
- May get stuck in local minima


## Learning rate too large:

- May overshoot minimum
- Can oscillate or diverge
- Learning becomes unstable

## Key Insight:

Finding the right learning rate is crucial for efficient learning.

# Learning Rate Comparison



- **Red path:** Learning rate too small ( $\eta = 0.1$ )
- **Blue path:** Good learning rate ( $\eta = 0.5$ )
- **Green path:** Learning rate too large ( $\eta = 1.5$ )