

# Weight Sharing in Convolutional Neural Networks

Sayan CHAKI

April 9, 2025

# Outline

- 1 What is Weight Sharing?
- 2 Why Do We Share Weights?
- 3 Visual Explanation
- 4 Mathematical Perspective
- 5 Intuitive Analogies
- 6 Thinking Deeper
- 7 Summary


# What is Weight Sharing?

- A fundamental concept in convolutional neural networks (CNNs)
- In traditional fully-connected networks: each input-output connection has **unique weights**
- In CNNs: the same weights are reused across different spatial locations

**The same filter weights are applied at every position in the input**

# The Convolutional Layer

- A **filter/kernel** (e.g.,  $3 \times 3$ ) slides across the input image
- At each position, performs the same computation
- Produces an output feature map
- The **same filter weights** are used at **every spatial location**



# Why Share Weights?

## Key Advantages

- **Translation Equivariance**
  - Objects are detected regardless of their position in the image
  - If we shift the input, the output shifts accordingly
- **Parameter Efficiency**
  - Drastically reduces the number of learnable parameters
  - Makes training feasible for large images
- **Better Generalization**
  - Captures spatial patterns regardless of location
  - Strong inductive bias for visual data

# Parameter Efficiency: A Comparison

Consider processing a  $32 \times 32$  image:

## Fully Connected Layer:


- Each output neuron connected to all 1,024 inputs
- For 64 output neurons:  
 $1,024 \times 64 = 65,536$  parameters

## Convolutional Layer:

- 8 filters of size  $3 \times 3$
- Only  $3 \times 3 \times 8 = 72$  parameters
- **910× fewer parameters!**

**Weight sharing enables deep architectures that can process high-resolution images with reasonable computational resources**

# Visual Explanation: Weight Sharing



**The same filter weights  $w_{ij}$   
are used at every position!**

# Mathematical Definition

For a filter  $W$  of size  $k \times k$  and an input image  $I$ :

$$\text{Output}(i, j) = \sum_{m=0}^{k-1} \sum_{n=0}^{k-1} W(m, n) \cdot I(i + m, j + n) \quad (1)$$

$$= (W * I)(i, j) \quad (2)$$

Where:

- $W(m, n)$  = Weight at position  $(m, n)$  in the filter
- $I(i + m, j + n)$  = Input value at offset position
- $*$  = Convolution operation

**Critical point:** The weights  $W(m, n)$  remain the **same** for all spatial positions  $(i, j)$  in the input image.



# Intuitive Analogies

## The Stencil Analogy

- A stencil pressed across different areas of a page
- The **same pattern** is applied everywhere
- Different input regions produce different outputs
- But the **stencil itself** doesn't change

## The Cookie Cutter Analogy

- One cookie cutter shapes multiple cookies
- **Reuses the same tool** across the dough
- Different dough regions (inputs) produce different cookies (outputs)
- But the **cutter itself** is identical

In both cases, we use the **same tool** across different locations to detect patterns or features.

# Food for Thought

## What if we didn't share weights?

- For a  $224 \times 224$  image with  $3 \times 3$  filters:
  - We'd need  $3 \times 3 \times 222 \times 222$  parameters per filter!
  - $\approx 148,000$  parameters vs. just 9 with weight sharing
- Training would be extremely inefficient
- Models would overfit drastically

## When might weight sharing not be ideal?

- When processing data with strong positional dependencies
- Example: Face recognition where features like eyes/nose have fixed positions
- Some architectures use partially locally-connected layers without full weight sharing

# Summary: Benefits of Weight Sharing

- **Drastically Reduces Parameters**

- Makes deep networks feasible
- Enables faster training

- **Provides Translation Equivariance**


- Detects patterns regardless of position

- **Improves Generalization**

- Strong inductive bias for spatial data

- **Enables Processing Large Images**

- Parameters independent of input size



**Weight sharing is what makes convolutional networks so powerful and efficient for visual tasks!**